

標題：HW3-report

姓名：葉怡君

學號：110062529

Implementation

Hw3-1

```
for(int k = 0; k < vertices; k++){
    #pragma omp parallel for num_threads(ncpus)
    for(int i = 0; i < vertices; i++){
        for(int j = 0; j < vertices; j++){
            if(board[i][j] > (board[i][k] + board[k][j]))
                board[i][j] = board[i][k] + board[k][j];
        }
    }
}
```

我選用了 FLOYD-WARSHALL ALGORITHM 去實作，只將 i 的那層 loop parallize。

Hw3-2

一開始先 cudaMalloc 一個 d_Dist 的陣列，在 input 的 function 用 unroll 的概念 read 至 Dist 以減少 I/O time，再用 cudaMemcpy 至 d_Dist，再根據 blocking factor(=32)去分 block，blocking factor = 32 是因為 warp size = 32，32*32 threads/block 是因為 blocking factor = 32，每個 grid 有多少個 block 取決於當時所在的 phase 及正在計算的 row 或 column，ex. 正在 phase 2，計算 pivot block 右邊的 row，則將 blocks/grid 設成 blockDim.x = round - r - 1 及 blockDim.y = 1，接著，每次 calculate 都會先把 data load 到 shared memory 去做計算，計算完成後再存回去 device 的 global memory，以減少 memory access 的時間，在每個 phase 結束前都用 cudaDeviceSynchronize() 同步，才可開始下個階段，最後 cudaMemcpy 至 Dist 做 output。

Hw3-3

在 hw3-3，我有試著使用 omp_thread 分給兩個 device 去跑，但發現並沒有跑得比較快，反而有些只用一個 gpu 跑、有 accepted 的測資，分給兩個 gpu 跑的時候卻超時，因此 hw3-3 的 code 跟 3-2 是一樣的。

Profiling(hw3-2)

實驗設備：課堂 hades

測資：c06.1

● occupancy

```
[pp21s49@hades01 ~]$ srun -p prof -N1 -n1 --gres=gpu:1 nvprof --metrics achieved_occupancy ./hw3-2 /home/pp21/share/hw3-2/cases/c06.1 tmp.out
==893021== NVPROF is profiling process 893021, command: ./hw3-2 /home/pp21/share/hw3-2/cases/c06.1 tmp.out
==893021== Profiling application: ./hw3-2 /home/pp21/share/hw3-2/cases/c06.1 tmp.out
devices count: 1
I/O time: 2.000000 ms
kernel time: 190.000000 ms
mem copy time: 0.000000 ms
==893021== Profiling result:
==893021== Metric result:
Invocations      Metric Name      Metric Description      Min      Max      Avg
Device "GeForce GTX 1080 (0)"
Kernel: cal_1(int, int, int, int*, int)
4      achieved_occupancy      Achieved Occupancy      0.304458  0.496400  0.448160
Kernel: cal_3(int, int, int, int*, int)
10     achieved_occupancy      Achieved Occupancy      0.316964  0.484857  0.434347
Kernel: cal_col(int, int, int, int*, int)
6      achieved_occupancy      Achieved Occupancy      0.137417  0.496208  0.397101
Kernel: cal_row(int, int, int, int*, int)
6      achieved_occupancy      Achieved Occupancy      0.315967  0.496649  0.466300
```

● sm efficiency

```
[pp21s49@hades01 ~]$ srun -p prof -N1 -n1 --gres=gpu:1 nvprof --metrics sm_efficiency ./hw3-2 /home/pp21/share/hw3-2/cases/c06.1 tmp.out
==893047== NVPROF is profiling process 893047, command: ./hw3-2 /home/pp21/share/hw3-2/cases/c06.1 tmp.out
==893047== Profiling application: ./hw3-2 /home/pp21/share/hw3-2/cases/c06.1 tmp.out
devices count: 1
I/O time: 2.000000 ms
kernel time: 128.000000 ms
mem copy time: 0.000000 ms
==893047== Profiling result:
==893047== Metric result:
Invocations
Device "GeForce GTX 1080 (0)"
Kernel: cal_1(int, int, int, int*, int)
4 sm_efficiency Multiprocessor Activity 1.74% 3.69% 3.20%
Kernel: cal_3(int, int, int, int*, int)
10 sm_efficiency Multiprocessor Activity 1.83% 19.79% 7.42%
Kernel: cal_col(int, int, int, int*, int)
6 sm_efficiency Multiprocessor Activity 3.24% 9.91% 6.10%
Kernel: cal_row(int, int, int, int*, int)
6 sm_efficiency Multiprocessor Activity 3.60% 11.23% 6.43%
```

● shared memory load/store throughput

```
[pp21s49@hades01 ~]$ srun -p prof -N1 -n1 --gres=gpu:1 nvprof --metrics shared_load_throughput,shared_store_throughput ./hw3-2 /home/pp21/share/hw3-2/cases/c06.1 tmp.out
==893072== NVPROF is profiling process 893072, command: ./hw3-2 /home/pp21/share/hw3-2/cases/c06.1 tmp.out
==893072== Profiling application: ./hw3-2 /home/pp21/share/hw3-2/cases/c06.1 tmp.out
devices count: 1
I/O time: 4.000000 ms
kernel time: 170.000000 ms
mem copy time: 0.000000 ms
==893072== Profiling result:
==893072== Metric result:
Invocations
Device "GeForce GTX 1080 (0)"
Kernel: cal_1(int, int, int, int*, int)
4 shared_load_throughput Shared Memory Load Throughput 3.0268GB/s 58.401GB/s 50.192GB/s
4 shared_store_throughput Shared Memory Store Throughput 1.8626GB/s 10.499GB/s 8.2010GB/s
Kernel: cal_3(int, int, int, int*, int)
10 shared_load_throughput Shared Memory Load Throughput 14.234GB/s 446.21GB/s 177.78GB/s
10 shared_store_throughput Shared Memory Store Throughput 2.5636GB/s 33.528GB/s 9.8766GB/s
Kernel: cal_col(int, int, int, int*, int)
6 shared_load_throughput Shared Memory Load Throughput 9.9684GB/s 115.06GB/s 75.351GB/s
6 shared_store_throughput Shared Memory Store Throughput 2.2095GB/s 10.638GB/s 7.6691GB/s
Kernel: cal_row(int, int, int, int*, int)
6 shared_load_throughput Shared Memory Load Throughput 8.8052GB/s 165.19GB/s 95.834GB/s
6 shared_store_throughput Shared Memory Store Throughput 2.8670GB/s 33.262GB/s 15.474GB/s
```

● global load/store throughput

```
[pp21s49@hades01 ~]$ srun -p prof -N1 -n1 --gres=gpu:1 nvprof --metrics gst_throughput ./hw3-2 /home/pp21/share/hw3-2/cases/c06.1 tmp.out
srun: job 420762 queued and waiting for resources
srun: job 420762 has been allocated resources
==727290== NVPROF is profiling process 727290, command: ./hw3-2 /home/pp21/share/hw3-2/cases/c06.1 tmp.out
devices count: 1
I/O time: 6.000000 ms
kernel time: 58.000000 ms
mem copy time: 0.000000 ms
==727290== Profiling application: ./hw3-2 /home/pp21/share/hw3-2/cases/c06.1 tmp.out
==727290== Profiling result:
==727290== Metric result:
Invocations
Device "GeForce GTX 1080 (0)"
Kernel: cal_1(int, int, int, int*, int)
4 gst_throughput Global Store Throughput 52.982MB/s 632.85MB/s 540.70MB/s
Kernel: cal_3(int, int, int, int*, int)
10 gst_throughput Global Store Throughput 50.863MB/s 17.243GB/s 3.2566GB/s
Kernel: cal_col(int, int, int, int*, int)
6 gst_throughput Global Store Throughput 109.34MB/s 1.2028GB/s 843.83MB/s
Kernel: cal_row(int, int, int, int*, int)
6 gst_throughput Global Store Throughput 131.54MB/s 1.2789GB/s 795.89MB/s
```

```
[pp21s49@hades01 ~]$ srun -p prof -N1 -n1 --gres=gpu:1 nvprof --metrics gld_throughput ./hw3-2 /home/pp21/share/hw3-2/cases/c06.1 tmp.out
==893638== NVPROF is profiling process 893638, command: ./hw3-2 /home/pp21/share/hw3-2/cases/c06.1 tmp.out
==893638== Profiling application: ./hw3-2 /home/pp21/share/hw3-2/cases/c06.1 tmp.out
devices count: 1
I/O time: 10.000000 ms
kernel time: 196.000000 ms
mem copy time: 0.000000 ms
==893638== Profiling result:
==893638== Metric result:
Invocations
Device "GeForce GTX 1080 (0)"
Kernel: cal_1(int, int, int, int*, int)
4 gld_throughput Global Load Throughput 290.34MB/s 468.70MB/s 334.33MB/s
Kernel: cal_3(int, int, int, int*, int)
10 gld_throughput Global Load Throughput 625.17MB/s 6.2106GB/s 2.3562GB/s
Kernel: cal_col(int, int, int, int*, int)
6 gld_throughput Global Load Throughput 456.98MB/s 1.7048GB/s 959.78MB/s
Kernel: cal_row(int, int, int, int*, int)
6 gld_throughput Global Load Throughput 443.71MB/s 1.7305GB/s 935.53MB/s
```

Experiment & Analysis

實驗設備：課堂 hades

測資：c2l.1

時間計算方式： `std::chrono::steady_clock`

Computing time：在執行 block_FW 所花的時間

Memory copy time：進行 `cudaMemcpy - HostToDevice & DeviceToHost` 所花的時間

I/O time：在 read 跟 write 檔案時所花的時間

- **Time distribution(hw3-2)：**

Computing time: 334.000000 ms

Memory copy time: 30.000000 ms

I/O time: 712.000000 ms

- **Optimization(hw3-2)**

在 input 函式中使用 `unroll` 的概念來存進 Dist。

使用 `blockIdx & threadIdx` 平行化(CUDA 2D alignment)得把需要用來判斷的 block 的 data load 至 shared memory，再利用 shared memory 中的 data 來計算 shortest path，計算完成後再存回 `Dist(device' s global variable)`。

Experience & Conclusion

這次作業我學到如何寫 cuda，以及更深的了解 GPU 的 `grid`、`block`、`thread` 的概念，還有 device 的 `global memory` 跟 `shared memory access` 的時間真的差很多，這次作業難度很高，在 cuda 的撰寫遇到了很多困難，在如何設置 `blockIdx` 和 `threadIdx` 這個地方就花了很多時間，即便做了 `blockIdx` 和 `threadIdx` 的平行化及使用 `shared memory`，在 hw3-2 p19 以後的測資還是沒過，最後在 hw3-3，我有試著使用 `omp_thread` 分給兩個 device 去跑，但發現並沒有跑得比較快，蠻想知道兩個 gpu 的寫法。